

# A Deep Learning Approach for Innovating MATSim Plans


Fred Shone - UCL BIG - [frederick.shone.17@ucl.ac.uk](mailto:frederick.shone.17@ucl.ac.uk)

Dr Tim Hillel - UCL BIG - [tim.hillel@ucl.ac.uk](mailto:tim.hillel@ucl.ac.uk)

# Contents

1. Context!
2. Problem Statement
3. Experiment A - Plan to Score Model
4. Experiment B - Plan to Plan Model


# Context - DGMs for Activity Sequence Generation



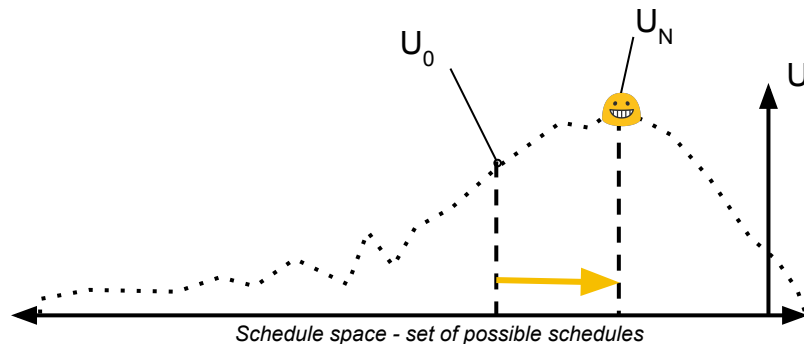
- Fast (on a GPU)
- “Simultaneous” choice generation
- Able to learn realism implicitly from data
- No need to specify structure or alternative specific variables

# Problem Statement

Consider all the possible utility scores an agent can experience in simulation, depending on their mode, route and time choices (*but ignoring other agents for simplicity*):



At iteration 0 we start at some initial plan with  $U_0$ .  
At some iterations an *innovation* or *innovations* take place, with which we hope to discover a better plan:




# Problem Statement

But what are the odds that our innovation strategies are proposing good plans?

Note:

Agents do not learn from bad experiences, so every simulated bad plan is a wasted iteration.

Bad plans also “poison/taint” the innovations made by other agents. Which is why we limit the overall rate of innovation/exploration.



# Quick Experiment

Existing MATSim model of Sheffield, UK. 120,000 agents. Active modes, routed cycling, transit inter-modal access-egress.

We consider the best utility of agents after intervals of 50 iterations. After each interval we report the probability that an agent finds a better plan. We can think of this as the *efficiency* of the strategies responsible for innovations.

Iterations	% agents find a better plan	...per innovation*	...per iteration
50 -> 100	73%	4.9%	1.5%
100 -> 150	62%	4.1%	1.2%
150 -> 200**	56%	3.7%	1.1%***

\* Innovation rate (probability of mode, route or time innovation) is 30%.

\*\* Final iterations include cool down/annealing.

\*\*\* Arguably this should be zero for final iterations as agents arrive at their optima and cannot do better.

# Innovation Strategies

Consider the approaches to plan innovations, ie the “*strategies*”:

Strategy Name	Type	Critique
ChangeExpBeta	Exploit	This is required to (i) update plan scores for new choices by other agents and (ii) create a realistic simulation.
RouteChoice	Shortest path	Edge costs based on previous iteration/s. Edge costs need to be complete/correct for each agent. Greedy.
ModeChoice	Random	Does not use previous experience by agent or other agents. No/limited consideration of relationship of other choices within plan.
TimeChoice	Random walk	
<i>DMC extension</i>	<i>Exploit</i>	<i>Mode innovations only. Requires calculation/estimation of alternatives/utilities. No consideration of relationship of other choices types within plan.</i>

# Implications

- Inefficient strategies for innovations = slow discovery of *more* optimal plans.
- Innovation rate has to be limited, else supply simulation is misleading.
- Time innovation is often removed/restricted.
- Hardcoded rules to reduce bad mode choices.
- Need high quality initial plans.
- Sim size limited.
- Can't reasonably add more complex choices to big sims.



# The Problem Statement

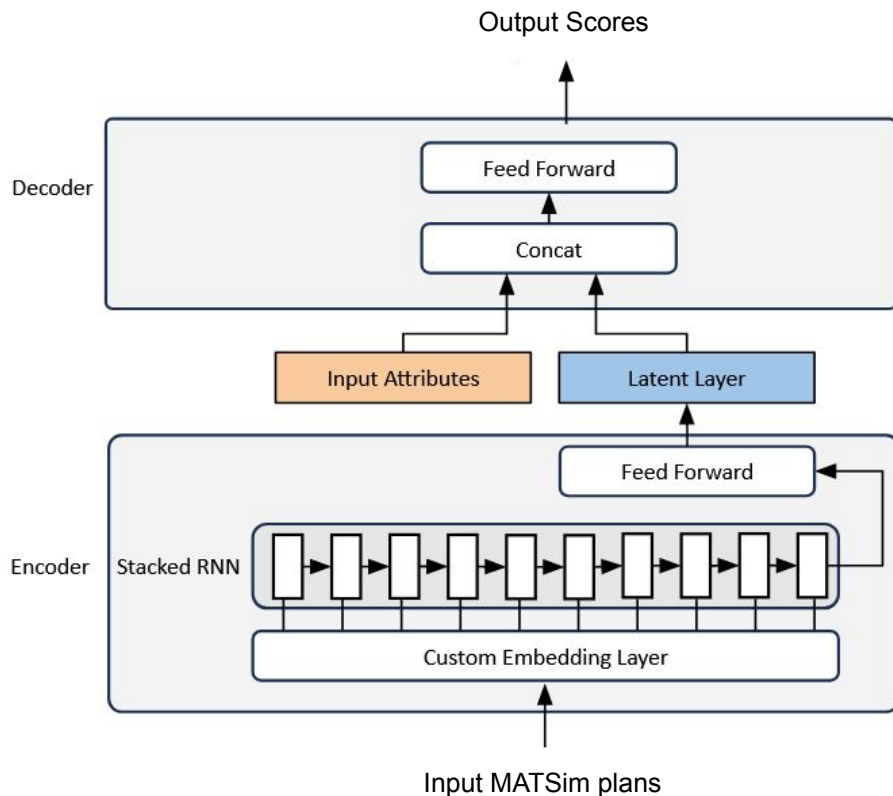
Can we approximate the distribution of agent utilities over their schedule spaces such that innovations can generate schedules that are more likely to have a higher utility?"

Can think of this as wanting to improve the “*efficiency*” of innovations.

Find new approach for **mode** and **time** innovations that:

- Considers combined choice(s) within a plan - “*simultaneous*”
- Considers experience by other agents - “*shared*”
- Considers experience from previous iterations - “*remembered*”

# Experiment A - Plan to Score Model



Plans are encoded as sequences of activities and trips, padded with special tokens to make the lengths consistent. Each component has a duration, mode and distance.

Attributes are one-hot encoded.

The encoder block uses a learnt embedding for the activity and mode types. Followed by stacked LSTM layers.

This intermediate “latent” layer is concatenated with the encoded attributes, then fed through a regression block of fully connected layers to output a utility score for each sequence. Mean squared error is used for loss.

Many more details at [github/fredshone/caveat](https://github.com/fredshone/caveat)

# Encoding MATSim Plans

(i) 24-hour long activity sequences are assembled from trip information. Composed of activity types and travel (categorical):

(ii) Extract durations, normalise with max duration:



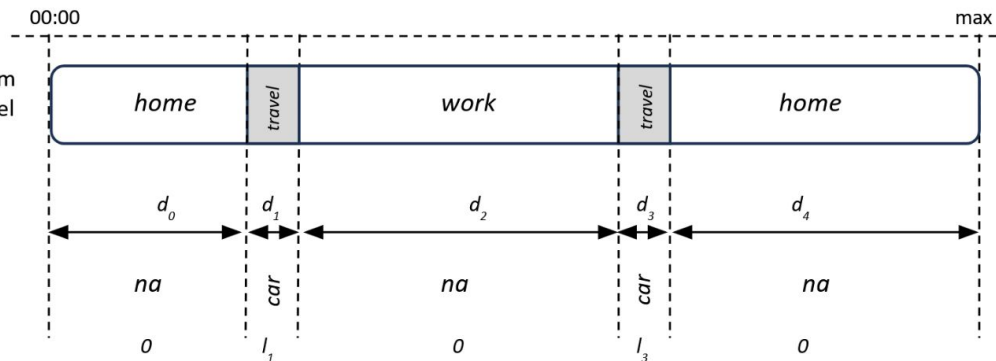
(iii) Extract modes (categorical):

(iv) Extract distances, normalise with max distance:

(v) Combine into vectors:

(vi) Sequence padded, up to some maximum length, with special start and end of sequence characters:


(vii) Final vector representation:



0	2	3	4	3	2	1	1	1	1
0	.2	.1	.3	.1	.3	0	0	0	0
0	0	1	0	1	0	0	0	0	0
0	0	.3	0	.3	0	0	0	0	0


# Training

- Extract plans and associated scores from iterations *50, 100, 150 and 200* of an existing MATSim run of 120,000 agents, ~0.5 million plans.
- Train, validate, test split of (80/10/10).
- Training loss (MSE) is **0.0084** (abs error is ~0.1 Utils, equivalent to 0.1GBP).
- Model is learning to approximate **both** the simulation and scoring function.



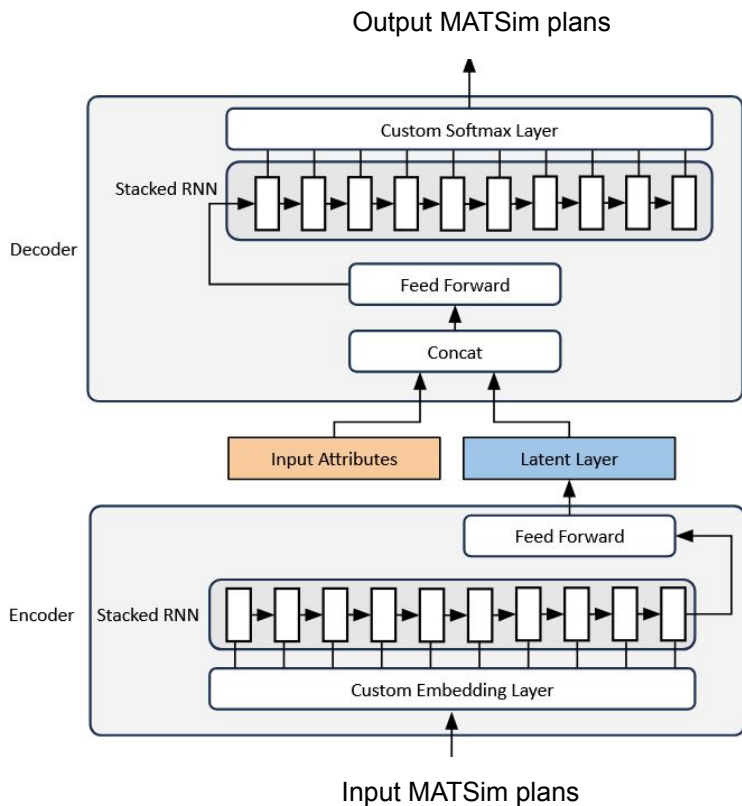
# Example Application/Discussion

1. Replanning module proposes N candidates
2. Model picks best/ignores worst
3. Lagged update model training



- **Improve quality of plans**
- Most additional compute is on GPU
- **Potential to transfer learn from previous runs**
- Easy to parameterise exploration/exploitation
- Memory implications for creating proposals
- Implementation pretty tough
- **Cold startup/training update**
- **Far simpler models might also work well**

# Experiment B - Plan to Plan Model



Plans are encoded as sequences of activities and trips, padded with special tokens to make the lengths consistent. Each component has a type, duration, mode and distance.

Attributes are one-hot encoded.


The encoder block uses a learnt embedding for the activity and mode types. Followed by stacked LSTM layers.

This intermediate “latent” layer is concatenated with the encoded attributes, then fed through a decoder block which is similar to the encoder. We use a custom combination of losses.


Many more details at [github/fredshone/caveat](https://github.com/fredshone/caveat)

# Training

- Extract plans and associated scores from iterations *50, 100, 150 and 200* of an existing MATSim run of 120,000 agents.
- Pair together plans for each agent, such that the LHS (input) plan has a lower utility than the RHS (target) plan. In total we create 375,000 pairs.



Training pairs:




# Training


- We train the model to directly infer a “*better plan*” for given input plan.
- Where “better” is a schedule with trip times and modes that will simulate with higher utility.
- We use a 80/10/10 split.
- The test set has an average plan improvement of **4.3** utils.
- We evaluate the utility of generated plans using the matsim scoring function, but note this is an estimate only, as plans have not been simulated.
- The model is able to generate plans with an average improvement of 2.2 utils.
- But on closer inspection, generates an improvement only 36% of the time.
- More pragmatic evaluation:
  - B-** Correctly identifies the target mode 35% of the time.
  - F** Correctly identifies that an activity should be shorter/longer 50% of the time.
- ~5% failure rate (inferring a schedule with altered activity sequence).




### The Good



### The Bad




### The Ugly



# Application/Discussion

1. Use model as an innovation strategy directly
2. Lagged update model training



- Improve quality of plans
  - **Additional work on GPU**
  - Potential to transfer learning from previous scenarios
  - **Can still parameterise exploration/exploitation by combining with other strategies**
- 
- Implementation pretty tough
  - **Cannot guarantee viable plans**
  - **Potential for failure to explore all choices**
  - Cold startup/training update

# Conclusions

- Demonstrated two approaches to improve the “quality” or “efficiency” of innovations for plan modes and times:
  - *Simultaneous* consideration of choices
  - *Shared* experience of choices
  - *Remembered* experience of choices
- Potentially (very?) fast, but some outstanding questions around implementation.
- Cost and risk high. But hope of interest for further research.
- Demonstrates the power and flexibility of deep learning approaches...
- Broadly we have trained faster (concurrent & differentiable) proxy models. If you like this work then there are other applications:
  - Activity schedule synthesis/generation
  - Influence on behavioural theory
  - Latent representations can be used to measure distances

# Further Work

- Activity schedule synthesis:
  - Synthetic datasets
  - Upsampling for simulations
  - Modelling as part of an activity-based model
- Increase model capabilities (trips, modes, multi-days).
- Identify and disseminate best performing encodings and architectures (CAVEAT).
- Apply to a multi-agent simulation optimisation problem - EV charging?

**Fred Shone** - frederick.shone.17@ucl.ac.uk

Dr Hillel - UCL BIG - tim.hillel@ucl.ac.uk

Dr Perez-Ortiz - UCL AI Centre

<https://fredshone.github.io/>

[frederick.shone.17@ucl.ac.uk](mailto:frederick.shone.17@ucl.ac.uk)

[github.com/fredshone](https://github.com/fredshone)

Thank you Arup City Modelling Group!

DCM extension - [github.com/matsim-org/matsim-lib/tree/master/contribs/discrete\\_mode\\_choice](https://github.com/matsim-org/matsim-lib/tree/master/contribs/discrete_mode_choice)

CAVEAT - [github/fredshone/caveat](https://github.com/fredshone/caveat)

PAM - [github/arup-group/pam](https://github.com/arup-group/pam)